Crowd Monitoring & QA

[Week 7]

This Week's Goals

- Finalize and thoroughly test the integrated pipeline and API, ensuring that individual video frames can be processed end-to-end to produce both crowd density heatmaps and JSON-based people counts, and that the system works seamlessly with the backend team for future integration.
- Collect, carefully annotate, and organize frames extracted from AFL footage to create a comprehensive custom dataset for YOLO, ensuring the data is well-prepared for the upcoming training and accurate evaluation of the crowd detection model.

1. **Build a pipeline that integrates both steps and create an API to enable collaboration with the backend team**

Assigned to: Son Tung, Jake

Outcomes:

- Successfully combined **LISA for Audience Segmentation** and **CSRNet for Crowd Density Estimation** into a single pipeline, enabling end-to-end processing of video frames.
- Developed an **API** that takes individual video frames as input and outputs (1) a heatmap image showing crowd density, and (2) a JSON file with the estimated number of people in the stage area.
- Wrote a **Markdown-based API usage guide** documenting input requirements, headers, request body, and expected responses, making it easier for the backend team to integrate and test the system.

Crowd Estimation API Usage Guide

This API processes individual video frames to generate:
1. A **heatmap image** showing crowd density.
2. A **JSON file** estimating the number of people in the stage area.

Endpoint
`POST /api/v1/crowd-estimation`

Headers
Key	Value
`Content-Type`	`multipart/form-data`
`Accept`	`application/json`

Input Requirements
- Input must be a **single video frame** (image file).
- Supported formats: `.jpg`, `.jpeg`, `.png`.
- Recommended resolution: 720p or higher for best accuracy.

Request Body (multipart/form-data)
Field	Type	Description
`frame`	File	The video frame image to be analyzed

Example (cURL):
```bash
curl -X POST http://localhost:8000/api/v1/crowd-estimation \
  -H "Accept: application/json" \
  -F "frame=@/path/to/frame.jpg"
```

2. **Train YOLO on frames extracted from AFL footage samples (2 weeks)**

Assigned to: Anika, Saksham, Binul

Outcomes this week:

- Test YOLOv8 for person detection on AFL snapshots with the custom dataset
- Collected and annotated frames extracted from AFL footage samples to prepare a custom training dataset for YOLO.

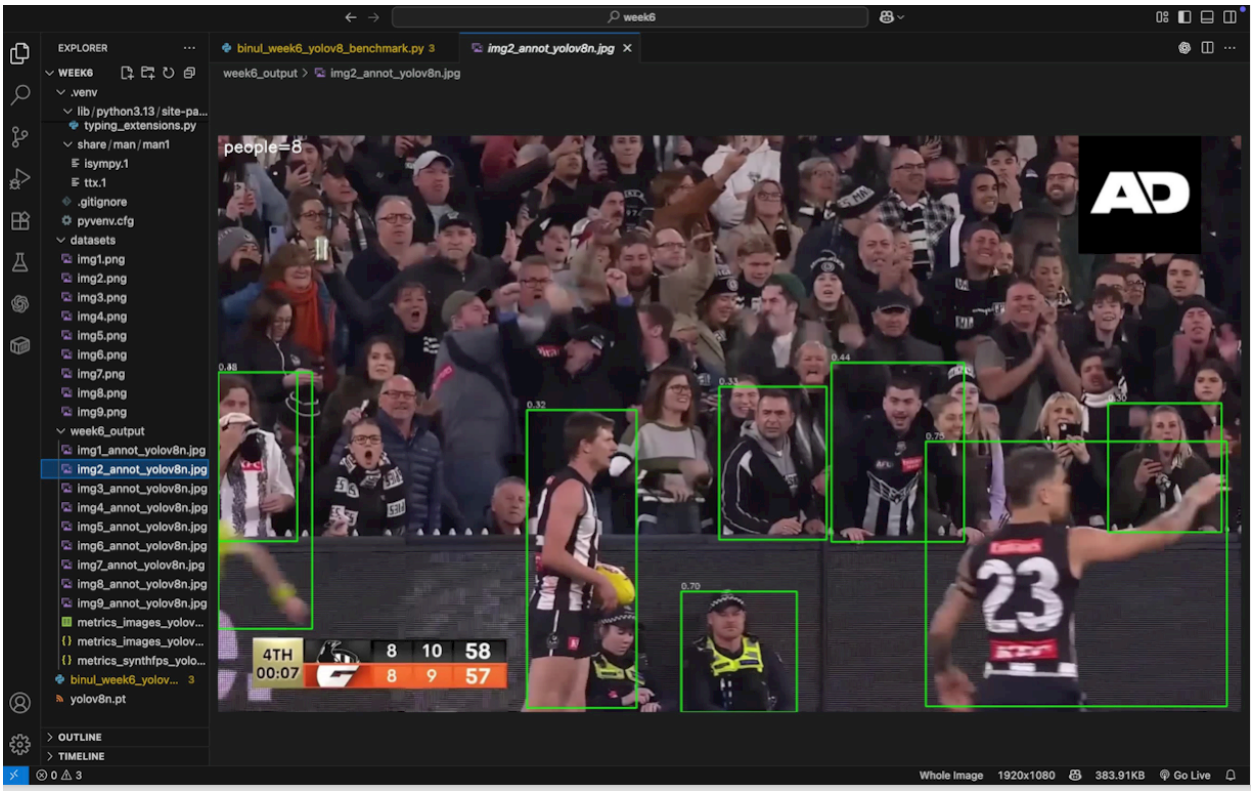

Note: The results were obtained using models with pretrained COCO weights, not AFL-specific weights.

Plan for next week:

- Trained YOLO on the annotated dataset to detect crowd regions within the frames and evaluated model accuracy using metrics such as mAP.
- Compared model performance before and after training to measure improvements in crowd detection capability.